

Write Up for P2: Advanced Lane Lines

Introduction:	1
Camera calibration	1
Image pipeline	1
Distortion correction	2
Binary Transform	2
Perspective Transform	3
Identify Lanes	3
Fit Lanes	4
Quality assessment:	4
Decision based on fit quality	4
Fit lanes to the image	5
Compute curvature	5
Mix project image + original	5

Introduction:

This document is a write up for the second project. As mentioned in the write up template, the objectives are:

- Compute the camera calibration matrix and distortion coefficients given a set of chessboard images.
- Apply a distortion correction to raw images.
- Use color transforms, gradients, etc., to create a thresholded binary image.
- Apply a perspective transform to rectify binary image ("birds-eye view").
- Detect lane pixels and fit to find the lane boundary.
- Determine the curvature of the lane and vehicle position with respect to center.
- Warp the detected lane boundaries back onto the original image.
- Output visual display of the lane boundaries and numerical estimation of lane curvature and vehicle position.

In order to make the report more consistent, I decided to make the description for code blocks that were already provided in the course brief and more detailed personal on code.

Project structure is as follows:

Code: **/P2.ipynb**

Videos: Input= **/test_videos** , Output= **/output_videos**

Images: Input= **/test_images** , Output= **/output_images**

Each cell in the pipeline is named and numbered, so code parts are referred to by cell name/number (+lines range if the cell contains more than one function). The code structure follows the rubrik steps: Camera calibration -> Image pipeline -> Video pipeline.

Camera calibration

I use the camera calibration function provided in the course without modification. The code is in script format (not a function) and takes as input n_x & n_y . Below before/after example is provided [here](#).

Code location: **First: camera calibration using chessboard images**

Image pipeline

The pipeline consists of one function: **image_pipeline** located in **4. Pipeline Function**

Dependencies (cell names):

- **Camera calibration parameters:** These are a one-time-cost, such as camera matrix and distortion coefficients.
- **1. Helper functions:** Contains all the helper functions developed for this project.
- **2. Source points for each road:** Used to perform perspective transform and contain the coordinates for 3 polygons, one for each video input.
- **3. Tracker class:** Contains project global variables, mostly useful for videos.

Side code:

- **5. Convert png images (from video) to jpg:** Converts video screenshots from .png to .jpg

- **6. Finding source points:** Manually get coordinates of lanes
- **7. Parameters tuning:** Self explanatory
- **8. Load and process image:** performs the pipeline on image then saves the result

Distortion correction

Function Name: None, it's one line of code: `cv2.undistort`

Location: Called directly within the **image_pipeline**, line 6

Input: Image + camera parameters

Output: undistorted image

Description: The before/after example provided below serves to illustrate camera calibration as well. Distortion correction is visible on the car's hood

Before



After



Binary Transform

Function name: **create_binary**

Location: **1. Helper functions**, lines 3->21

Functions called inside: **None**

Input: Image, l_channel threshold, b_channel threshold


Output: Binary image

Description: After trying several color spaces and gradient methods, I found that LAB colorspace provides the best overall performance. **L**, **B** channels to detect white, yellow colors respectively. I discarded gradients method because I found that they were very prone to noise which disturbed the lane finding algorithm.

My approach was to favor cleanness of prediction with a risk of an empty lane to robustness with a risk of high noise. I have to be clear though, image processing is a craft that I don't muster! My approach was empirical and not based on prior knowledge or experience but I'm thankful that I learned a bit in this project.

Before

After



Perspective Transform

Function name: transform_perspective

Location: 1.Helper Functions, Line: 27 -> 36

Functions called inside: 2. Source points for each road

Input: image, source_point_coordinates, offset


output:warped_img, M, Minv

Description: The code for this function is identical to the course: Use a manually picked coordinate as input and “offsetted” polygon in the target image.


It's important to report that offset + coordinates have to be adapted to the image otherwise the results may be very different, I report 2 reasons:

1. Lanes location can be different from one road to the other. That's why I set 3 different coordinates according to the road.
2. A large offset implies fitting the polynome to a larger surface which means more wiggling in the fit, on the other hand, a small offset does not capture road curvatures

Before



After



Identify Lanes

Function name: find_lane_lines

Location: 1.Helper Functions, Line: 39 -> 45

Functions called inside:

-lanes_from_previous, located: 1.Helper Functions, Line: 47 -> 64

-sliding_window, located: 1.Helper Functions, Line: 66 -> 127

Input: image,nwindows,margin,minpix,margin2

Output: leftx,lefty,rightx,righty,left_found,right_found

Description: This function identifies lanes in the input image but does not fit. If lanes in previous image were found then I use **lanes_from_previous**, if not **sliding_window**. The output consists of x,y coordinates for left/right lanes + 2 binary variables: Lane pixels found or not. Almost all code is provided in the course, I have added 2 restrictions:

-If one lane is 5 times longer than the other \Rightarrow the short one is considered not found.

-If length of the longest lane $< 1/2$ image height \Rightarrow both are considered not found.

The reason is that most often, the polynomial is badly fitted if the length is too short.

Fit Lanes

Function name: **fit_lane_lines**

Location: **1.Helper Functions, Line: 130 -> 306**

Functions called inside:

-Line class, located: **3. Tracker class**

-compute_curvature, located: **1.Helper Functions, Line: 309 -> 315**

Input:

-image,

-leftx, lefty, rightx, righty, left_found, right_found, #output from previous lane finding function

-smoothing (Use smoothing or not), xm=3.7/700 (width pixel resolution), ym=30/720 (height pixel resolution)

Output: image with fitted lanes

Description: The function can be divided into 3 distinct parts: Assess quality of found lanes \rightarrow Decide what to do based on the quality \rightarrow fit to the image either by smoothing or not. I'll go through each step.

NB: One key assumption I use is road width constance.

Quality assessment:

Location: **Line 142 \rightarrow 205**

Input: Image, lefty, rightx, righty, left_found, right_found

Output: 2 booleans: **good_left_fit & good_right_fit**

Description:

We define an error margin named "alpha" (around 30%) for the assessment of quality.

1. For each lane: Initialize to bad fit

If we found pixels \rightarrow Fit polynomial \rightarrow If previous lane was fitted \rightarrow compare curvature + quadratic coefficient + intercept with previous \rightarrow If OK for all \rightarrow fit is good

2. For both lanes and if previous tests were successful:

Compare quadratic coefficients \rightarrow If NOK \rightarrow Shorter lane is no good fit

Decision based on fit quality

Location: **Line 209 \rightarrow 289**

Input: **good_left_fit & good_right_fit**

Output (approximate): fitted lanes + updated global coefficients

Description: The decision part is 2 folds: Find the best fit + update global variables

1. If both fits are bad:

Fitting: Last best fit stored in the global variable

- Update: None
2. If one fit is bad:

Fitting:

 - If previous frame was well fitted for that lane → Use coefficients from previous frame
 - If not → Use current good frame coefficients + compute the road width → Adjust intercept according to the computed width

Update:

Road curvature from the good lane + polynomial fits for both lanes + set the bad lane to not found
 3. If both fits are good:

Fitting: Intuitive, Update: Analogous to previous


Fit lanes to the image

Location: Line 290 → 304

Input (approximate): Decided fits for current frame

Output (approximate): Fitted lanes to the image

Description: By now we have the fits resulting from the current frame. Depending on the decision to smooth frames or not + size of the smoother we output the final frames.



Compute curvature

Function name: `compute_curvature`

Location: 1.Helper Functions, Line: 309 -> 315

Functions called inside: None

Input: `imY, ypix, xpix, xm, ym`

Output: Radius

Description: The code for this block is identical to the one provided in the course.

Mix project image + original

Function name: `weighted_img`

Location: 1.Helper Functions, Line: 318 -> 322

Functions called inside: None

Input: original image, projected image

Output: final image

Description: I add the radius + car drift during this step



Video pipeline

Project video & challenge video link:

https://github.com/mmarouen/CarND-Advanced-Lane-Lines-master/tree/master/output_videos

Harder challenge link:

<https://drive.google.com/open?id=1B2gh8qizghYTNrWQ9u3e7G3uIWzUVHlu>

Discussion

A word about difficulty

Besides the method, the list of parameters by step that impact final result includes:

Binary image: Thresholds

Perspective transform: Offset amount + coordinates of the source polygon

Lane finding & fitting: nwindows, margin, minpix, smoothing size, error sensitivity

Image processing is a craft in itself, the challenge for me was to connect the dots and come up with a "reasonable" result given my current knowledge.

Improvement areas

1. The harder challenge made me aware of potential improvement possibilities:

-When one lane is not in the image for high curvature

-Detect/Process high light images

-Detect double lanes and find the center

- Deal with high curvatures: When the polynomial fit crosses the vertical half of the image
- 2. Fixing smoother size inversely proportional to car speed